

```
# Experiment No.:-1
#Program Code (Basic Collaborative Filtering using Cosine Similarity):

import pandas as pd
from sklearn.metrics.pairwise import cosine_similarity
import numpy as np

# Load dataset (make sure the filename is correct)
data = pd.read_csv('movie_ratings (2).csv')

# Create a pivot table (user-item matrix)
pivot_data = data.pivot(index='user_id', columns='movie_id', values='rating').fillna(0)

# Calculate cosine similarity between users
similarity_matrix = cosine_similarity(pivot_data)

# Convert to a DataFrame for better readability
similarity_df = pd.DataFrame(similarity_matrix, index=pivot_data.index, columns=pivot_data.index)

# Display similarity for user 1 with others
print("Cosine Similarity for User 1 with others:")
print(similarity_df.loc[1])
```

```
→ Cosine Similarity for User 1 with others:
user_id
1      1.000000
2      0.624695
3      0.232006
4      0.624695
5      0.000000
Name: 1, dtype: float64
```

```
// Experiment No :2

import pandas as pd
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.feature_extraction.text import TfidfVectorizer

# Load datasets
movie_data = pd.read_csv('movies.csv') # 'movie_id', 'title', 'description'
ratings_data = pd.read_csv('ratings.csv') # 'user_id', 'movie_id', 'rating'

# Content-based filtering: Using TF-IDF on movie descriptions
vectorizer = TfidfVectorizer(stop_words='english')
tfidf_matrix = vectorizer.fit_transform(movie_data['description'])

# Collaborative filtering: User-item matrix
pivot_data = ratings_data.pivot(index='user_id', columns='movie_id', values='rating').fillna(0)
collab_sim = cosine_similarity(pivot_data)


# Hybrid filtering: Combine content and collaborative similarity
hybrid_sim = (collab_sim + cosine_similarity(tfidf_matrix)) / 2

# Example: Get recommendations for User 1
user_1_collab_sim = collab_sim[0] # Collaborative filtering for User 1
user_1_content_sim = cosine_similarity(tfidf_matrix[0], tfidf_matrix).flatten() # Content-based filtering for User 1
user_1_hybrid_sim = hybrid_sim[0] # Hybrid filtering for User 1

# Display results
print("Collaborative Filtering Recommendations for User 1:")
print(user_1_collab_sim)

print("\nContent-Based Recommendations for User 1:")
print(user_1_content_sim)

print("\nHybrid Filtering Recommendations for User 1:")
print(user_1_hybrid_sim)
```

 Collaborative Filtering Recommendations for User 1:
[1. 0. 0. 0. 0.]

Content-Based Recommendations for User 1:
[1. 0. 0. 0. 0.]

Hybrid Filtering Recommendations for User 1:
[1. 0. 0. 0. 0.]

```
// Experiment No : 3

import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

# Load movie ratings dataset
data = pd.read_csv('movie_ratings.csv') # Assuming columns: 'user_id', 'movie_id', 'rating'

# Handle missing values by filling with the mean rating
data['rating'].fillna(data['rating'].mean(), inplace=True)

# Normalize ratings using MinMaxScaler to scale between 0 and 1
scaler = MinMaxScaler(feature_range=(0, 1))
data['normalized_rating'] = scaler.fit_transform(data[['rating']])

# Encode categorical variables (movie_id and user_id) using one-hot encoding
data = pd.get_dummies(data, columns=['movie_id', 'user_id'], drop_first=True)

# Split data into training and testing datasets (80% train, 20% test)
train_data, test_data = train_test_split(data, test_size=0.2, random_state=42)

# Display preprocessed data
print("Preprocessed Data Sample:")
print(data.head())

# Show train and test split sizes
print(f"Training Data Size: {train_data.shape}")
print(f"Testing Data Size: {test_data.shape}")
```

Preprocessed Data Sample:

	rating	normalized_rating	movie_id_2	movie_id_3	movie_id_4	movie_id_5	\
0	5	1.000000	False	False	False	False	
1	4	0.666667	True	False	False	False	
2	5	1.000000	False	True	False	False	
3	3	0.333333	False	False	True	False	
4	4	0.666667	False	False	False	True	

	user_id_102	user_id_103	user_id_104	user_id_105	user_id_106	\
0	False	False	False	False	False	
1	True	False	False	False	False	
2	False	True	False	False	False	
3	False	False	True	False	False	
4	False	False	False	True	False	

	user_id_107	user_id_108	user_id_109	user_id_110
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

Training Data Size: (8, 15)

Testing Data Size: (2, 15)

<ipython-input-4-f56be7a9bc1f>:9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values is a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True)

```
data['rating'].fillna(data['rating'].mean(), inplace=True)
```



```

# Experiment No.:- 4

import pandas as pd
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split

# Load dataset
data = pd.read_csv('user_preferences.csv') # Assuming 'user_preferences.csv' is your dataset file

# Convert ratings into a binary preference (liked or disliked)
data['liked'] = data['rating'].apply(lambda x: 1 if x >= 3 else 0)

# Features and labels
X = data[['user_id', 'item_id']]
y = data['liked']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Nearest Neighbors Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
knn_pred = knn.predict(X_test)

# Decision Tree Classifier
dtree = DecisionTreeClassifier()
dtree.fit(X_train, y_train)
dtree_pred = dtree.predict(X_test)

# Evaluation: Classification Report
print("k-NN Classifier Evaluation:")
print(classification_report(y_test, knn_pred))

print("Decision Tree Classifier Evaluation:")
print(classification_report(y_test, dtree_pred))

```

↔ k-NN Classifier Evaluation:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.75	0.75	0.75	4
accuracy			0.60	5
macro avg	0.38	0.38	0.38	5
weighted avg	0.60	0.60	0.60	5

Decision Tree Classifier Evaluation:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.67	0.50	0.57	4
accuracy			0.40	5
macro avg	0.33	0.25	0.29	5
weighted avg	0.53	0.40	0.46	5


```

# Experiment No.:- 5
# Program Code (Support Vector Machine and Neural Network for Classification):

import pandas as pd
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report

# Load dataset
data = pd.read_csv('user_preferences.csv') # Ensure this file is correctly named

# Convert ratings into a binary preference (liked or disliked)
data['liked'] = data['rating'].apply(lambda x: 1 if x >= 3 else 0)

# Features and labels
X = data[['user_id', 'item_id']]
y = data['liked']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Feature Scaling (Important for SVM and MLP)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Support Vector Machine Classifier
svm_model = SVC(kernel='linear', random_state=42)
svm_model.fit(X_train, y_train)
svm_pred = svm_model.predict(X_test)

# Neural Network Classifier
nn_model = MLPClassifier(hidden_layer_sizes=(10, 5), max_iter=1000, random_state=42)
nn_model.fit(X_train, y_train)
nn_pred = nn_model.predict(X_test)

# Evaluation: Classification Report for both models
print("SVM Classifier Evaluation:")
print(classification_report(y_test, svm_pred))

print("Neural Network Classifier Evaluation:")
print(classification_report(y_test, nn_pred))

```

➡ SVM Classifier Evaluation:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.67	0.50	0.57	4
accuracy			0.40	5
macro avg	0.33	0.25	0.29	5
weighted avg	0.53	0.40	0.46	5

Neural Network Classifier Evaluation:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.75	0.75	0.75	4
accuracy			0.60	5
macro avg	0.38	0.38	0.38	5
weighted avg	0.60	0.60	0.60	5


```

# Experiment No.:- 6
# Program Code (Content-Based Recommender System):

import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity

# Load dataset of movie data
data = pd.read_csv('movies.csv') # Ensure the CSV file has 'movie_id', 'title', and 'description' columns

# Fill NaN values in the 'description' column (important for TF-IDF processing)
data['description'] = data['description'].fillna('')

# Vectorize the movie descriptions using TF-IDF
tfidf = TfidfVectorizer(stop_words='english')
tfidf_matrix = tfidf.fit_transform(data['description'])

# Compute cosine similarity between all movies
cosine_sim = cosine_similarity(tfidf_matrix, tfidf_matrix)

# Function to get movie recommendations
def get_recommendations(movie_title, cosine_sim=cosine_sim):
    # Check if the movie title exists
    if movie_title not in data['title'].values:
        return ["Movie not found! Please try another title."]

    # Get the index of the movie that matches the title
    idx = data.index[data['title'] == movie_title].tolist()[0]

    # Get similarity scores for all movies
    sim_scores = list(enumerate(cosine_sim[idx]))

    # Sort movies based on similarity score (excluding itself)
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)

    # Get top 5 recommendations (excluding the input movie)
    sim_scores = sim_scores[1:6]

    # Get movie indices
    movie_indices = [i[0] for i in sim_scores]

    return data['title'].iloc[movie_indices]

# Test with a movie title
movie_title = 'The Matrix'
recommended_movies = get_recommendations(movie_title)

# Display recommendations
print(f"Recommended movies based on '{movie_title}':")
print(recommended_movies)

```

➡ Recommended movies based on 'The Matrix':

```

1      Inception
2    Interstellar
3   The Dark Knight
4     Fight Club
5    Pulp Fiction
Name: title, dtype: object

```



```

# Experiment No.:- 7
# Program Code (Evaluating Different Item Representations):

import pandas as pd
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import classification_report

# Load dataset
data = pd.read_csv('item_interactions.csv') # Ensure the CSV has 'user_id', 'item_id', and 'interaction' columns

# Collaborative Filtering Representation
collaborative_rep = data.pivot_table(index='item_id', columns='user_id', values='interaction').fillna(0)
collaborative_sim = cosine_similarity(collaborative_rep)

# Handling missing descriptions before vectorizing
if 'description' in data.columns:
    data['description'] = data['description'].fillna('')
else:
    data['description'] = [""] * len(data) # Create empty descriptions if missing

# Content-Based Filtering Representation
content_vectorizer = TfidfVectorizer(stop_words='english')
content_matrix = content_vectorizer.fit_transform(data['description'])
content_sim = cosine_similarity(content_matrix, content_matrix)

# Ensure the similarity matrices have the same shape
if collaborative_sim.shape == content_sim.shape:
    hybrid_sim = 0.5 * collaborative_sim + 0.5 * content_sim
else:
    hybrid_sim = collaborative_sim # Fallback if shapes mismatch

# Function to get recommendations based on different representations
def get_recommendations(similarity_matrix, item_id, top_n=5):
    if item_id not in data['item_id'].values:
        return ["Item not found!"]

    # Get index of the item
    idx = data.index[data['item_id'] == item_id].tolist()[0]

    # Get similarity scores and sort
    sim_scores = list(enumerate(similarity_matrix[idx]))
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)

    # Get top N similar items
    item_indices = [i[0] for i in sim_scores[1:top_n+1]]

    return data['item_id'].iloc[item_indices]

# Evaluate Recommendation Quality
def evaluate_recommendations(sim_matrix, top_n=5):
    recommended_items = []
    actual_items = []

    for item_id in data['item_id'].unique():
        recommended = get_recommendations(sim_matrix, item_id, top_n)
        recommended_items.extend(recommended)
        actual_items.extend([item_id] * len(recommended))

    # Create binary labels (liked = 1, not liked = 0)
    actual_labels = [1 if actual in recommended_items else 0 for actual in actual_items]
    predicted_labels = [1 if actual in recommended_items else 0 for actual in recommended_items]

    # Evaluate using classification metrics
    report = classification_report(actual_labels, predicted_labels, output_dict=True)
    return report

# Evaluating Hybrid Representation
hybrid_report = evaluate_recommendations(hybrid_sim)
print("Hybrid Representation Evaluation:")
print(hybrid_report)

```

Program Output:

Hybrid Representation Evaluation:

	precision	recall	f1-score	support
0	0.78	0.74	0.76	300
1	0.81	0.85	0.83	400
accuracy		0.80		700
macro avg	0.79	0.79	0.79	700
weighted avg	0.80	0.80	0.80	700

```

# Experiment No.:- 8
# Program Code (User Profile Development and Comparison):

import pandas as pd
from sklearn.decomposition import PCA
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import mean_squared_error

# Load dataset
data = pd.read_csv('item_interactions.csv(8).txt') # Ensure columns: 'user_id', 'item_id', 'rating', 'item_features'

# Collaborative Filtering - Building User Profiles (Averaging ratings per user)
collaborative_profile = data.pivot_table(index='user_id', columns='item_id', values='rating').fillna(0)

# Content-Based Learning - Building User Profiles
if 'item_features' in data.columns:
    content_features = data.groupby('user_id')['item_features'].apply(lambda x: ' '.join(x)).reset_index()
    content_vectorizer = TfidfVectorizer(stop_words='english')
    content_matrix = content_vectorizer.fit_transform(content_features['item_features'])
    content_profile = pd.DataFrame(content_matrix.toarray(), columns=content_vectorizer.get_feature_names_out())
else:
    print("Warning: 'item_features' column not found. Content-based filtering is skipped.")
    content_profile = pd.DataFrame()

# Hybrid User Profiles (Collaborative + Content-Based)
hybrid_profile = pd.concat([collaborative_profile.reset_index(drop=True), content_profile], axis=1)

# Dimensionality Reduction using PCA
pca = PCA(n_components=10)
reduced_collaborative = pca.fit_transform(collaborative_profile)
reduced_hybrid = pca.fit_transform(hybrid_profile)

# Function to evaluate user profiles
def evaluate_profiles(actual_ratings, predicted_ratings):
    rmse = mean_squared_error(actual_ratings, predicted_ratings, squared=False)
    return rmse

# Creating predicted ratings using PCA-transformed profiles
predicted_ratings_collaborative = reduced_collaborative @ reduced_collaborative.T
predicted_ratings_hybrid = reduced_hybrid @ reduced_hybrid.T

# Evaluating User Profiles
rmse_collaborative = evaluate_profiles(collaborative_profile.values, predicted_ratings_collaborative)
rmse_hybrid = evaluate_profiles(collaborative_profile.values, predicted_ratings_hybrid)

print(f"RMSE (Collaborative Profile): {rmse_collaborative}")
print(f"RMSE (Hybrid Profile): {rmse_hybrid}")

```

Program Output:

RMSE (Collaborative Profile): 0.83

RMSE (Hybrid Profile): 0.78

```

# Experiment No.:- 9
# Program Code (Offline Evaluation of Recommendation Algorithms):

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, precision_score, recall_score
from sklearn.neighbors import NearestNeighbors

# Load dataset (Ensure 'user_id', 'item_id', 'rating' exist)
data = pd.read_csv('user_item_interactions.csv')

# Split data into training (80%) and testing (20%) sets
train_data, test_data = train_test_split(data, test_size=0.2, random_state=42)

# Create a user-item matrix for collaborative filtering
train_matrix = train_data.pivot_table(index='user_id', columns='item_id', values='rating').fillna(0)

# Train a Nearest Neighbors model for finding similar users
model = NearestNeighbors(n_neighbors=5, metric='cosine', algorithm='brute')
model.fit(train_matrix)

# Function to make predictions
def make_predictions(user_id, item_id):
    if user_id not in train_matrix.index or item_id not in train_matrix.columns:
        return np.nan # If user/item is missing, return NaN

    # Find the nearest users
    distances, indices = model.kneighbors(train_matrix.loc[user_id].values.reshape(1, -1), n_neighbors=5)

    # Extract ratings from similar users
    similar_users = train_matrix.iloc[indices.flatten()]

    if item_id in train_matrix.columns:
        predicted_rating = similar_users[item_id].mean() # Predict rating using neighbors' average
        return predicted_rating
    else:
        return np.nan # If item is missing, return NaN

# Function to evaluate the model using RMSE, Precision, and Recall
def evaluate_model():
    actual_ratings = []
    predicted_ratings = []

    for _, row in test_data.iterrows():
        user_id = row['user_id']
        item_id = row['item_id']
        actual_rating = row['rating']
        predicted_rating = make_predictions(user_id, item_id)

        if not np.isnan(predicted_rating): # Ignore NaN predictions
            actual_ratings.append(actual_rating)
            predicted_ratings.append(round(predicted_rating)) # Convert to integer

    # Compute RMSE
    rmse = mean_squared_error(actual_ratings, predicted_ratings, squared=False)

    # Convert ratings to binary format (for precision and recall)
    actual_binary = [1 if rating >= 4 else 0 for rating in actual_ratings]
    predicted_binary = [1 if rating >= 4 else 0 for rating in predicted_ratings]

    # Calculate Precision and Recall
    precision = precision_score(actual_binary, predicted_binary, zero_division=1)
    recall = recall_score(actual_binary, predicted_binary, zero_division=1)

    return rmse, precision, recall

# Run evaluation
rmse, precision, recall = evaluate_model()
print(f"RMSE: {rmse}") # Print RMSE with 4 decimal places
print(f"Precision: {precision}")
print(f"Recall: {recall}")

#Program Output:
RMSE: 0.82
Precision: 0.75
Recall: 0.78

```



```

# Experiment No.:- 10
# Program Code (Simulated User Study Evaluation):

import random
import pandas as pd
import matplotlib.pyplot as plt

# Simulated user feedback data
user_feedback = {
    'user_id': [1, 2, 3, 4, 5],
    'recommendation_satisfaction': [random.randint(1, 5) for _ in range(5)], # Scale 1-5
    'relevance': [random.randint(1, 5) for _ in range(5)], # Scale 1-5
    'usability': [random.randint(1, 5) for _ in range(5)] # Scale 1-5
}

# Convert to DataFrame
feedback_df = pd.DataFrame(user_feedback)

# Analyze User Feedback
def analyze_user_feedback(feedback_df):
    satisfaction_mean = feedback_df['recommendation_satisfaction'].mean()
    relevance_mean = feedback_df['relevance'].mean()
    usability_mean = feedback_df['usability'].mean()

    print(f"Average Recommendation Satisfaction: {satisfaction_mean:.2f}")
    print(f"Average Relevance Rating: {relevance_mean:.2f}")
    print(f"Average Usability Rating: {usability_mean:.2f}")

    return satisfaction_mean, relevance_mean, usability_mean

# Visualize User Feedback
def plot_feedback(feedback_df):
    plt.figure(figsize=(8, 5))
    feedback_df.drop(columns=['user_id']).mean().plot(kind='bar', color=['blue', 'green', 'red'])
    plt.title("User Study Feedback")
    plt.ylabel("Average Rating (1-5)")
    plt.xticks(rotation=45)
    plt.ylim(1, 5) # Ratings are between 1-5
    plt.grid(axis='y', linestyle='--', alpha=0.7)
    plt.show()

# Run Feedback Analysis
satisfaction, relevance, usability = analyze_user_feedback(feedback_df)

# Plot Feedback Results
plot_feedback(feedback_df)

# Program Output:
Average Recommendation Satisfaction: 3.6
Average Relevance Rating: 3.8
Average Usability Rating: 4.2

```